

SHI FENG

shifeng@fas.harvard.edu | fengshi.link | [Google Scholar](#)

EDUCATION

Harvard University *Sept 2023 - Present*
Ph.D. in Computer Science, advised by [Prof. Yiling Chen](#)

Harvard University *Sept 2023 - Feb 2026*
S.M. in Computer Science, GPA: 3.93/4.0

Tsinghua University *Aug 2019 - Jun 2023*
B.E. (Honors) in Computer Science, selected to [Yao Class](#), GPA: 3.89/4.0

RESEARCH FOCUS

Large Language Models: post-training, reasoning, and benchmarks.

Economics and Computation: algorithmic game theory, market design, auctions, and strategic learning.

I develop reliable post-training and decision-making methods for strategic, noisy, or interactive data.

APPOINTMENTS

Graduate Research Assistant, Harvard University, advised by [Prof. Yiling Chen](#) *Sept 2023 - Present*

Research Intern, Microsoft Research Asia Theory Center, mentored by [Prof. Wei Chen](#) *Dec 2022 - Jun 2023*

Research Intern, EconCS Group, Harvard University, mentored by [Prof. Yiling Chen](#) *Feb 2022 - Aug 2022*

Research Intern, Microsoft Research Asia Theory Center, mentored by [Prof. Wei Chen](#) *Nov 2020 - Jan 2022*

WORKING PAPERS

(* denotes equal contribution or alphabetical order)

Yiling Chen*, **Shi Feng***, Sadie Zhao*.

Contextual Procurement Auctions with Bandit Learning.

In submission to **AAAI 2027**. [[arXiv](#)]

Shi Feng, Tuo Zhao, Yiling Chen.

When and What to Price Against Model Extraction.

In submission to **ICLR 2027**.

Shi Feng, Yiling Chen, Sham Kakade, Hanlin Zhang.

MILO: Mode-wise Likelihood Optimization for Correct-Mode Coverage.

In submission to **ICLR 2027**.

Hanlin Zhang, **Shi Feng**, David Parkes, Yiling Chen, Sham Kakade.

PRISMA: Dynamic Benchmarking as Mechanism Design.

In submission to **ICLR 2027**.

Shi Feng, Siwei Wang, Shang-Hua Teng, Wei Chen.

Mechanisms of Learned Invalid-Transition Detection in Autoregressive Planning.

In submission to **ICLR 2027**.

PUBLICATIONS

Yiling Chen*, **Shi Feng***, Grant Schoenebeck*, Fang-Yi Yu*.

Information Elicitation by Ordinal Agreement.

International Conference on Web and Internet Economics (WINE 2026).

Shi Feng*, Hanlin Zhang*, Fan Nie, Sham Kakade, Yiling Chen.

Peer-Predictive Self-Training for Language Model Reasoning.

Conference on Language Modeling (COLM 2026). [\[arXiv\]](#)

Yiling Chen*, **Shi Feng***, Paul Kattuman*, Fang-Yi Yu*.

Data Reliability Scoring.

The Twenty-Seventh ACM Conference on Economics and Computation (EC 2026). [\[arXiv\]](#)

Russell Li, **Shi Feng**, Safwan Hossain, Yiling Chen.

Strategic Fine-Tuning: Learning from Human Feedback under Incentive Misalignment.

Incentive-Based AI Alignment Workshop, EC 2026. [\[PDF\]](#)

Siwei Wang*, Yifei Shen*, Haoran Sun*, **Shi Feng***, Shang-Hua Teng, Li Dong, Yaru Hao, Wei Chen.

Benefits and Pitfalls of Reinforcement Learning for Language Model Planning: A Theoretical Perspective.

The Fourteenth International Conference on Learning Representations (ICLR 2026). [\[arXiv\]](#)

Yiling Chen*, **Shi Feng***, Fang-Yi Yu*.

Carrot and Stick: Eliciting Comparison Data and Beyond.

The 38th Conference on Neural Information Processing Systems (NeurIPS 2024). [\[arXiv\]](#)

Siwei Wang*, Yifei Shen*, **Shi Feng**, Haoran Sun, Shang-Hua Teng, Wei Chen.

ALPINE: Unveiling The Planning Capability of Autoregressive Learning in Language Models.

The 38th Conference on Neural Information Processing Systems (NeurIPS 2024). [\[arXiv\]](#)

Shi Feng*, Nuoya Xiong*, Wei Chen.

Combinatorial Causal Bandits without Graph Skeleton.

The 16th Asian Conference on Machine Learning (ACML 2024). [\[arXiv\]](#)

Shi Feng, Wei Chen.

Combinatorial Causal Bandits.

The 37th AAAI Conference on Artificial Intelligence (AAAI 2023). [\[arXiv\]](#)

Shi Feng, Fang-Yi Yu, Yiling Chen.

Peer Prediction for Learning Agents.

The 36th Conference on Neural Information Processing Systems (NeurIPS 2022). [\[arXiv\]](#)

Shi Feng, Wei Chen.

Causal Inference for Influence Propagation – Identifiability of the Independent Cascade Model.

The 10th International Conference on Computational Data and Social Networks (CSoNet 2021), Best Paper. [\[arXiv\]](#)

SELECTED HONORS & AWARDS

Outstanding Undergraduate

Jun 2023, THU IIS

“Star of Tomorrow” Award of Excellence

Mar 2022, MSR Asia

Best Paper Award (1/57)

Nov 2021, CSonNet

“Challenge Cup” Science and Technology Competition, Top Grade Prize (<i>top 3%</i>)	<i>Apr 2021, THU</i>
Tsinghua Xuetaang Scholarship	<i>Aug 2019, THU</i>
Russian Mathematical Olympiad, Silver Medal	<i>Apr 2019, MSE of Russia</i>
Chinese Mathematical Olympiad, Gold Medal	<i>Nov 2018, CMS</i>

TALKS

ESIF-AIML 2026 , Data Reliability Scoring	<i>Jun 2026</i>
Harvard EconCS Seminar , Data Reliability Scoring	<i>Nov 2025</i>
ACML 2024 , Combinatorial Causal Bandits without Graph Skeleton	<i>Dec 2024</i>
EC-24 Workshop , Carrot and Stick: Eliciting Comparison Data and Beyond	<i>Jul 2024</i>
Harvard EconCS Seminar , Reveal Planning Capability of Autoregressive Learning in LLMs	<i>Mar 2024</i>
MSRA Theory Seminar , LLMs: Current Theoretical Analysis and Future Directions	<i>Mar 2023</i>
MSRA Online Paper Sharing , Combinatorial Causal Bandits	<i>Mar 2023</i>
AAAI 2023 , Combinatorial Causal Bandits	<i>Feb 2023</i>
NeurIPS 2022 , Peer Prediction for Learning Agents	<i>Oct 2022</i>
Yao Class Seminar , Peer Prediction for Learning Agents	<i>Sept 2022</i>
CSoNet 2021 , Causal Inference for Influence Propagation–Identifiability of the IC Model	<i>Nov 2021</i>

SERVICES

Journal Reviewer: Journal of Combinatorial Optimization (JOCO), IEEE Journal on Selected Areas in Information Theory (JSAIT), Transactions on Machine Learning Research (TMLR)

Conference Reviewer: ICML 2023/2024/2025/2026, NeurIPS 2023/2024/2025/2026, ICLR 2024/2025/2026, AAAI 2026/2027, ACML 2025/2027, WINE 2026

Teaching Fellow: Network Science: Theory and Algorithms (2023 Spring) in Tsinghua University, Economics and Computation (2025 Spring), Convex Optimization and Applications in Machine Learning (2026 Spring) in Harvard University