

SHI FENG

shifeng@fas.harvard.edu | [fengshi.link](#) | [Google Scholar](#)

EDUCATION

Harvard University	Sept 2023 - Present
Ph.D. in Computer Science, advised by Prof. Yiling Chen	
Harvard University	Sept 2023 - Feb 2026
S.M. in Computer Science, GPA: 3.93/4.0	
Tsinghua University	Aug 2019 - Jun 2023
B.E. (Honors) in Computer Science, selected to Yao Class , GPA: 3.89/4.0	

RESEARCH FOCUS

Large Language Models: post-training, reasoning, and benchmarks.
Economics and Computation: algorithmic game theory, market design, auctions, and strategic learning.
I develop reliable post-training and decision-making methods for strategic, noisy, or interactive data.

APPOINTMENTS

Graduate Research Assistant, Harvard University, advised by Prof. Yiling Chen	Sept 2023 - Present
Research Intern, Microsoft Research Asia Theory Center, mentored by Prof. Wei Chen	Dec 2022 - Jun 2023
Research Intern, EconCS Group, Harvard University, mentored by Prof. Yiling Chen	Feb 2022 - Aug 2022
Research Intern, Microsoft Research Asia Theory Center, mentored by Prof. Wei Chen	Nov 2020 - Jan 2022

WORKING PAPERS

(* denotes equal contribution or alphabetical order)

Yiling Chen*, **Shi Feng***, Sadie Zhao*.

Contextual Procurement Auctions with Bandit Learning.

In submission to **AAAI 2027**. [\[arXiv\]](#)

Shi Feng, Tuo Zhao, Sadie Zhao, Yiling Chen.

When and What to Price Against Model Extraction.

In submission to **ICLR 2027**.

Shi Feng, Yiling Chen, Sham Kakade, Hanlin Zhang.

MILO: Mode-wise Likelihood Optimization for Correct-Mode Coverage.

In submission to **ICLR 2027**.

Hanlin Zhang, **Shi Feng**, Anton Tsitsulin, Vahab Mirrokni, Ariel Procaccia, Yiling Chen, David Parkes, Sham Kakade.

PRISMA: Dynamic Benchmarking as Mechanism Design.

In submission to **ICLR 2027**.

Shi Feng, Siwei Wang, Shang-Hua Teng, Wei Chen.

Internal Detection of Invalid Transitions in Transformers.

In submission to **ICLR 2027**.

PUBLICATIONS

Yiling Chen*, **Shi Feng***, Grant Schoenebeck*, Fang-Yi Yu*.

Information Elicitation by Ordinal Agreement.

International Conference on Web and Internet Economics (WINE 2026).

Shi Feng*, Hanlin Zhang*, Fan Nie, Sham Kakade, Yiling Chen.

Peer-Predictive Self-Training for Language Model Reasoning.

Conference on Language Modeling (COLM 2026). [\[arXiv\]](#)

Yiling Chen*, **Shi Feng***, Paul Kattuman*, Fang-Yi Yu*.

Data Reliability Scoring.

The Twenty-Seventh ACM Conference on Economics and Computation (EC 2026). [\[arXiv\]](#)

Russell Li, **Shi Feng**, Safwan Hossain, Yiling Chen.

Strategic Fine-Tuning: Learning from Human Feedback under Incentive Misalignment.

Incentive-Based AI Alignment Workshop, EC 2026. [\[PDF\]](#)

Siwei Wang*, Yifei Shen*, Haoran Sun*, **Shi Feng***, Shang-Hua Teng, Li Dong, Yaru Hao, Wei Chen.

Benefits and Pitfalls of Reinforcement Learning for Language Model Planning: A Theoretical Perspective.

The Fourteenth International Conference on Learning Representations (ICLR 2026). [\[arXiv\]](#)

Yiling Chen*, **Shi Feng***, Fang-Yi Yu*.

Carrot and Stick: Eliciting Comparison Data and Beyond.

The 38th Conference on Neural Information Processing Systems (NeurIPS 2024). [\[arXiv\]](#)

Siwei Wang*, Yifei Shen*, **Shi Feng**, Haoran Sun, Shang-Hua Teng, Wei Chen.

ALPINE: Unveiling The Planning Capability of Autoregressive Learning in Language Models.

The 38th Conference on Neural Information Processing Systems (NeurIPS 2024). [\[arXiv\]](#)

Shi Feng*, Nuoya Xiong*, Wei Chen.

Combinatorial Causal Bandits without Graph Skeleton.

The 16th Asian Conference on Machine Learning (ACML 2024). [\[arXiv\]](#)

Shi Feng, Wei Chen.

Combinatorial Causal Bandits.

The 37th AAAI Conference on Artificial Intelligence (AAAI 2023). [\[arXiv\]](#)

Shi Feng, Fang-Yi Yu, Yiling Chen.

Peer Prediction for Learning Agents.

The 36th Conference on Neural Information Processing Systems (NeurIPS 2022). [\[arXiv\]](#)

Shi Feng, Wei Chen.

Causal Inference for Influence Propagation – Identifiability of the Independent Cascade Model.

The 10th International Conference on Computational Data and Social Networks (CSoNet 2021), Best Paper. [\[arXiv\]](#)

SELECTED HONORS & AWARDS

Outstanding Undergraduate	Jun 2023, THU IIS
“Star of Tomorrow” Award of Excellence	Mar 2022, MSR Asia
Best Paper Award (1/57)	Nov 2021, CSonNet
“Challenge Cup” Science and Technology Competition, Top Grade Prize (top 3%)	Apr 2021, THU
Tsinghua Xuetang Scholarship	Aug 2019, THU
Russian Mathematical Olympiad, Silver Medal	Apr 2019, MSE of Russia
Chinese Mathematical Olympiad, Gold Medal	Nov 2018, CMS

TALKS

ESIF-AIML 2026, Data Reliability Scoring	Jun 2026
Harvard EconCS Seminar, Data Reliability Scoring	Nov 2025
ACML 2024, Combinatorial Causal Bandits without Graph Skeleton	Dec 2024
EC-24 Workshop, Carrot and Stick: Eliciting Comparison Data and Beyond	Jul 2024
Harvard EconCS Seminar, Reveal Planning Capability of Autoregressive Learning in LLMs	Mar 2024
MSRA Theory Seminar, LLMs: Current Theoretical Analysis and Future Directions	Mar 2023
MSRA Online Paper Sharing, Combinatorial Causal Bandits	Mar 2023
AAAI 2023, Combinatorial Causal Bandits	Feb 2023
NeurIPS 2022, Peer Prediction for Learning Agents	Oct 2022
Yao Class Seminar, Peer Prediction for Learning Agents	Sept 2022
CSonNet 2021, Causal Inference for Influence Propagation–Identifiability of the IC Model	Nov 2021

SERVICES

Journal Reviewer: Journal of Combinatorial Optimization (JOCO), IEEE Journal on Selected Areas in Information Theory (JSAIT), Transactions on Machine Learning Research (TMLR)

Conference Reviewer: ICML 2023/2024/2025/2026, NeurIPS 2023/2024/2025/2026, ICLR 2024/2025/2026, AAAI 2026/2027, ACML 2025/2027, WINE 2026

Teaching Fellow: Network Science: Theory and Algorithms (2023 Spring) in Tsinghua University, Economics and Computation (2025 Spring), Convex Optimization and Applications in Machine Learning (2026 Spring) in Harvard University